

4.6 — Fixed-width integers and size_t

👤 ALEX 🕒 OCTOBER 8, 2021

In the previous lessons on integers, we covered that C++ only guarantees that integer variables will have a minimum size -- but they could be larger, depending on the target system.

Why isn't the size of the integer variables fixed?

The short answer is that this goes back to C, when computers were slow and performance was of the utmost concern. C opted to intentionally leave the size of an integer open so that the compiler implementers could pick a size for `int` that performs best on the target computer architecture.

Doesn't this suck?

By modern standards, yes. As a programmer, it's a little ridiculous to have to deal with types that have uncertain ranges.

Consider the `int` type. The minimum size for `int` is 2 bytes, but it's often 4 bytes on modern architectures. If you assume an `int` is 4 bytes because that's most likely, then your program will probably misbehave on architectures where `int` is actually 2 bytes (since you will probably be storing values that require 4 bytes in a 2 byte variable, which will cause overflow or undefined behavior). If you assume an `int` is only 2 bytes to ensure maximum compatibility, then on systems where `int` is 4 bytes, you're wasting 2 bytes per integer and doubling your memory usage!

Fixed-width integers

To address the above issues, C99 defined a set of fixed-width integers (in the `stdint.h` header) that are guaranteed to be the same size on any architecture.

These are defined as follows:

Name	Type	Range	Notes
<code>std::int8_t</code>	1 byte signed	-128 to 127	Treated like a signed char on r
<code>std::uint8_t</code>	1 byte unsigned	0 to 255	Treated like an unsigned char
<code>std::int16_t</code>	2 byte signed	-32,768 to 32,767	
<code>std::uint16_t</code>	2 byte unsigned	0 to 65,535	
<code>std::int32_t</code>	4 byte signed	-2,147,483,648 to 2,147,483,647	
<code>std::uint32_t</code>	4 byte unsigned	0 to 4,294,967,295	
<code>std::int64_t</code>	8 byte signed	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807	
<code>std::uint64_t</code>	8 byte unsigned	0 to 18,446,744,073,709,551,615	

C++ officially adopted these fixed-width integers as part of C++11. They can be accessed by including the `<cstdint>` header, where they are defined inside the `std` namespace. Here's an example:

```

1  #include <cstdint> // for fixed-width
   integers
   #include <iostream>

   int main()
2  {
   std::int16_t i{5};
3  std::cout << i;
4  return 0;
5  }
```

The fixed-width integers have two downsides that are typically raised.

First, the fixed-width integers are not guaranteed to be defined on all architectures. They only exist on systems where there are fundamental types matching their widths and following a certain binary representation. Your program will fail to compile on any such architecture that does not support a fixed-width integer that your program is using. However, given that most modern architectures have standardized around 8/16/32/64-bit variables, this is unlikely to be a problem unless your program needs to be portable to some exotic mainframe or embedded architectures.

Second, if you use a fixed-width integer, it may be slower than a wider type on some architectures. For example, if you need an integer that is guaranteed to be 32-bits, you might decide to use `std::int32_t`, but your CPU might actually be faster at processing 64-bit integers. However, just because your CPU can process a given type faster doesn't mean your program will be faster overall -- modern programs are often constrained by memory usage rather than CPU, and the larger memory footprint may slow your program more than the faster CPU processing accelerates it. It's hard to know without actually measuring.

Fast and least integers

To help address the above downsides, C++ also defines two alternative sets of integers that are guaranteed to be defined.

The fast types (`std::int_fast#_t` and `std::uint_fast#_t`) provide the fastest signed/unsigned integer type with a width of at least # bits (where # = 8, 16, 32, or 64). For example, `std::int_fast32_t` will give you the fastest signed integer type that's at least 32 bits.

The least types (`std::int_least#_t` and `std::uint_least#_t`) provide the smallest signed/unsigned integer type with a width of at least # bits (where # = 8, 16, 32, or 64). For example, `std::uint_least32_t` will give you the smallest unsigned integer type that's at least 32 bits.

Here's an example from the author's Visual Studio (32-bit console application):

```

1 #include <cstdint> // for fixed-width integers
  #include <iostream>
2
3 int main()
4 {
5     std::cout << "least 8:  " << sizeof(std::int_least8_t) * 8 << "
6     bits\n";
7     std::cout << "least 16: " << sizeof(std::int_least16_t) * 8 << "
8     bits\n";
9     std::cout << "least 32: " << sizeof(std::int_least32_t) * 8 << "
10    bits\n";
11    std::cout << '\n';
12    std::cout << "fast 8:  " << sizeof(std::int_fast8_t) * 8 << " bits\n";
13    std::cout << "fast 16: " << sizeof(std::int_fast16_t) * 8 << "
14    bits\n";
15    std::cout << "fast 32: " << sizeof(std::int_fast32_t) * 8 << "
16    bits\n";
17
18    return 0;
19 }

```

This produced the result:

```

least 8:  8 bits
least 16: 16 bits
least 32: 32 bits

fast 8:  8 bits
fast 16: 32 bits
fast 32: 32 bits

```

You can see that `std::int_least16_t` is 16 bits, whereas `std::int_fast16_t` is actually 32 bits. This is because on the author's machine, 32-bit integers are faster to process than 16-bit integers.

However, these fast and least integers have their own downsides: First, not many programmers actually use them, and a lack of familiarity can lead to errors. Second, the fast types can lead to the same kind of memory wastage that we saw with 4 byte integers. Most seriously, because the size of the fast/least integers can vary, it's possible that your program may exhibit different behaviors on architectures where they resolve to different sizes. For example:

```

1 #include <cstdint> // for fixed-width integers
  #include <iostream>
2
3 int main()
4 {
5     std::uint_fast16_t sometype { 0 };
6     --sometype; // intentionally overflow to invoke wraparound
7     behavior
8
9     std::cout << sometype;
10
11    return 0;
12 }

```

This code will produce different results depending on whether `std::uint_fast16_t` is 16, 32, or 64 bits.

It's hard to know where your program might not function as expected until you've rigorously tested your program on such

architectures. And we imagine not many developers have access to a wide range of different architectures to test with!

Warning: `std::int8_t` and `std::uint8_t` may behave like chars instead of integers

Note: We talk more about chars in lesson [4.11 -- Chars](#)).

Due to an oversight in the C++ specification, most compilers define and treat `std::int8_t` and `std::uint8_t` (and the corresponding fast and least fixed-width types) identically to types *signed char* and *unsigned char* respectively. Consequently, `std::cin` and `std::cout` may work differently than you're expecting. Here's a sample program showing this:

```
1 #include <stdint>
2 #include <iostream>
3
4 int main()
5 {
6     std::int8_t
7     myint{65};
8     std::cout << myint;
9
10    return 0;
11 }
```

On most systems, this program will print 'A' (treating *myint* as a char). However, on some systems, this may print *65* as expected.

For simplicity, it's best to avoid `std::int8_t` and `std::uint8_t` (and the related fast and least types) altogether (use `std::int16_t` or `std::uint16_t` instead). However, if you do use `std::int8_t` or `std::uint8_t`, you should be careful of anything that would interpret `std::int8_t` or `std::uint8_t` as a char instead of an integer (this includes `std::cout` and `std::cin`).

Warning

Avoid the 8-bit fixed-width integer types. If you do use them, note that they are often treated like chars.

Integral best practices

Given the various pros and cons of the fundamental integral types, the fixed-width integral types, the fast/least integral types, and signed/unsigned challenges, there is little consensus on integral best practices.

Our stance is that it's better to be correct than fast, better to fail at compile time than runtime -- therefore, we recommend avoiding the fast/least types in favor of the fixed-width types. If you later discover the need to support a platform for which the fixed-width types won't compile, then you can decide how to migrate your program (and thoroughly test) at that point.

Best practice

- Prefer `int` when the size of the integer doesn't matter (e.g. the number will always fit within the range of a 2-byte signed integer). For example, if you're asking the user to enter their age, or counting from 1 to 10, it doesn't matter whether `int` is 16 or 32 bits (the numbers will fit either way). This will cover the vast majority of the cases you're likely to run across.
- Prefer `std::int#_t` when storing a quantity that needs a guaranteed range.
- Prefer `std::uint#_t` when doing bit manipulation or where well-defined wrap-around behavior is required.

Avoid the following when possible:

- Unsigned types for holding quantities
- The 8-bit fixed-width integer types
- The fast and least fixed-width types

- Any compiler-specific fixed-width integers -- for example, Visual Studio defines `_int8`, `_int16`, etc...

What is `std::size_t`?

Consider the following code:

```
1 #include <iostream>
2 int main()
3 {
4     std::cout << sizeof(int) <<
5     '\n';
6     return 0;
7 }
```

On the author's machine, this prints:

```
4
```

Pretty simple, right? We can infer that operator `sizeof` returns an integer value -- but what integer type is that return value? An `int`? A `short`? The answer is that `sizeof` (and many functions that return a size or length value) return a value of type `std::size_t`. `std::size_t` is defined as an unsigned integral type, and it is typically used to represent the size or length of objects.

Amusingly, we can use the `sizeof` operator (which returns a value of type `std::size_t`) to ask for the size of `std::size_t` itself:

```
1 #include <cstdint> // std::size_t
2 #include <iostream>
3 int main()
4 {
5     std::cout << sizeof(std::size_t) <<
6     '\n';
7     return 0;
8 }
```

Compiled as a 32-bit (4 byte) console app on the author's system, this prints:

```
4
```

Much like an integer can vary in size depending on the system, `std::size_t` also varies in size. `std::size_t` is guaranteed to be unsigned and at least 16 bits, but on most systems will be equivalent to the address-width of the application. That is, for 32-bit applications, `std::size_t` will typically be a 32-bit unsigned integer, and for a 64-bit application, `size_t` will typically be a 64-bit unsigned integer. `size_t` is defined to be big enough to hold the size of the largest object creatable on your system (in bytes). For example, if `std::size_t` is 4 bytes wide, the largest object creatable on your system can't be larger than 4,294,967,295 bytes, because this is the largest number a 4 byte unsigned integer can store. This is only the uppermost limit of an object's size, the real size limit can be lower depending on the compiler you're using.

By definition, any object with a size (in bytes) larger than the largest integral value `size_t` can hold is considered ill-formed (and will cause a compile error), as the `sizeof` operator would not be able to return the size without wrapping around.



Next lesson

4.7 Introduction to scientific notation



Back to table of contents



Previous lesson

4.5 Unsigned integers, and why to avoid them

Leave a comment... Put C++ code between triple-backticks (markdown style):````Your C++ code````

 Name*

 Email* 

Avatars from <https://gravatar.com/> are connected to your provided email address.

Notify me about replies:



POST COMMENT

DP N N FOUT

Newest ▼

