

4.12 — An introduction to std::string

ALEX AUGUST 29, 2021

The very first C++ program you wrote probably looked something like this:

```
1 #include <iostream>
2 int main()
3 {
4     std::cout << "Hello,
5     world!\n";
6     return 0;
7 }
```

So what is “Hello, world!” exactly? “Hello, world!” is a collection of sequential characters called a **string**. In C++, we use strings to represent text such as names, addresses, words, and sentences. String literals (such as “Hello, world!\n”) are placed between double quotes to identify them as strings.

Because strings are commonly used in programs, most modern programming languages include a fundamental string data type. In C++, strings aren’t a fundamental type (they’re actually a **compound type**, and defined in the C++ standard library rather than as part of the core language). But strings are straightforward and useful enough that we’ll introduce them here rather than wait until the chapter on compound types ([chapter 9](#)).

std::string

To use strings in C++, we first need to #include the <string> header to bring in the declarations for std::string. Once that is done, we can define variables of type std::string.

```
1 #include <string> // allows use of
   std::string
   std::string myName {}; // empty string
```

Just like normal variables, you can initialize or assign values to strings as you would expect:

```
1 | std::string myName{ "Alex" }; // initialize myName with string literal
   | "Alex"
   | myName = "John"; // assign variable myName the string literal "John"
```

Note that strings can hold numbers as well:

```
1 | std::string myID{ "45" }; // "45" is not the same as integer
   | 45!
```

In string form, numbers are treated as text, not numbers, and thus they can not be manipulated as numbers (e.g. you can't multiply them). C++ will not automatically convert string numbers to integer or floating point values.

String output

Strings can be output as expected using `std::cout`:

```
1 | #include <iostream>
   | #include <string>
2 |
3 | int main()
4 | {
5 |     std::string myName{ "Alex" };
6 |     std::cout << "My name is: " << myName <<
   | '\n';
   |
   |     return 0;
7 | }
```

This prints:

```
My name is: Alex
```

Empty strings will print nothing:

```
1 | #include <iostream>
   | #include <string>
2 |
3 | int main()
4 | {
5 |     std::string empty{ };
6 |     std::cout << '[' << empty <<
   | ']';
   |
   |     return 0;
7 | }
```

Which prints:

```
[ ]
```

String input with `std::cin`

Using strings with `std::cin` may yield some surprises! Consider the following example:

```
1 #include <iostream>
2 #include <string>
3
4 int main()
5 {
6     std::cout << "Enter your full name: ";
7     std::string name{};
8     std::cin >> name; // this won't work as expected since std::cin breaks on
    whitespace
9
10    std::cout << "Enter your age: ";
11    std::string age{};
12    std::cin >> age;
13
14    std::cout << "Your name is " << name << " and your age is " << age << '\n';
15
16    return 0;
17 }
```

Here's the results from a sample run of this program:

```
Enter your full name: John Doe
Enter your age: Your name is John and your age is Doe
```

Hmmm, that isn't right! What happened? It turns out that when using `operator>>` to extract a string from `cin`, `operator>>` only returns characters up to the first whitespace it encounters. Any other characters are left inside `std::cin`, waiting for the next extraction.

So when we used `operator>>` to extract a string into variable `name`, only `"John"` was extracted, leaving `" Doe"` inside `std::cin`. When we then used `operator>>` to get variable `age`, it extracted `"Doe"` instead of waiting for us to input an age. Then the program ends.

Use `std::getline()` to input text

To read a full line of input into a string, you're better off using the `std::getline()` function instead. `std::getline()` takes two parameters: the first is `std::cin`, and the second is your string variable.

Here's the same program as above using `std::getline()`:

```

1  #include <string> // For std::string and std::getline
   #include <iostream>

2  int main()
3  {
4      std::cout << "Enter your full name: ";
5      std::string name{};
6      std::getline(std::cin >> std::ws, name); // read a full line of text into
   name
7
8      std::cout << "Enter your age: ";
9      std::string age{};
10     std::getline(std::cin >> std::ws, age); // read a full line of text into
   age
11
12     std::cout << "Your name is " << name << " and your age is " << age << '\n';
13
14     return 0;
15 }

```

Now our program works as expected:

```

Enter your full name: John Doe
Enter your age: 23
Your name is John Doe and your age is 23

```

What the heck is std::ws?

In lesson [4.8 -- Floating point numbers](#), we discussed `output manipulators`, which allow us to alter the way output is displayed. In that lesson, we used the `output manipulator` function `std::setprecision()` to change the number of digits of precision that `std::cout` displayed.

C++ also supports input manipulators, which alter the way that input is accepted. The `std::ws` `input manipulator` tells `std::cin` to ignore any leading whitespace.

Let's explore why this is useful. Consider the following program:

```

1  #include <string>
2  #include <iostream>
3
4  int main()
5  {
6      std::cout << "Pick 1 or 2: ";
7      int choice{};
8      std::cin >> choice;
9
10     std::cout << "Now enter your name: ";
11     std::string name{};
12     std::getline(std::cin, name); // note: no std::ws here
13
14     std::cout << "Hello, " << name << ", you picked " << choice <<
   '\n';
15
16     return 0;
17 }

```

Here's some output from this program:

```

Pick 1 or 2: 2
Now enter your name: Hello, , you picked 2

```

This program first asks you to enter 1 or 2, and waits for you to do so. All good so far. Then it will ask you to enter your name. However, it won't actually wait for you to enter your name! Instead, it prints the "Hello" string, and then exits. What happened?

It turns out, when you enter a value using operator `>>`, `std::cin` not only captures the value, it also captures the newline character `'\n'` that occurs when you hit the `enter` key. So when we type `2` and then hit `enter`, `std::cin` gets the string `"2\n"`. It then extracts the `2` to variable `choice`, leaving the newline character behind for later. Then, when `std::getline()` goes to read the name, it sees `"\n"` is already in the stream, and figures we must have previously entered an empty string! Definitely not what was intended.

We can amend the above program to use the `std::ws` input manipulator, to tell `std::getline()` to ignore any leading whitespace characters:

```
1 #include <string>
2 #include <iostream>
3
4 int main()
5 {
6     std::cout << "Pick 1 or 2: ";
7     int choice{};
8     std::cin >> choice;
9
10    std::cout << "Now enter your name: ";
11    std::string name{};
12    std::getline(std::cin >> std::ws, name); // note: added std::ws
    here
13
14    std::cout << "Hello, " << name << ", you picked " << choice <<
15    '\n';
16
17    return 0;
18 }
```

Now this program will function as intended.

```
Pick 1 or 2: 2
Now enter your name: Alex
Hello, Alex, you picked 2
```

Best practice

If using `std::getline` to read strings, use the `std::ws` input manipulator to ignore leading whitespace.

Key insight

Using the extraction operator (`>>`) with `std::cin` ignores leading whitespace.
`std::getline` does not ignore leading whitespace unless you use input manipulator `std::ws`.

String length

If we want to know how many characters are in a `std::string`, we can ask the `std::string` for its length. The syntax for doing this is different than you've seen before, but is pretty straightforward:

```
1 #include <iostream>
2 #include <string>
3
4 int main()
5 {
6     std::string myName{ "Alex" };
7     std::cout << myName << " has " << myName.length() << "
8     characters\n";
9     return 0;
10 }
```

This prints:

```
Alex has 4 characters
```

Note that instead of asking for the string length as `length(myName)`, we say `myName.length()`. The `length()` function isn't a normal standalone function -- it's a special type of function that is nested within `std::string` called a `member function`. Because `length()` lives within `std::string`, it is sometimes written as `std::string::length` in documentation.

We'll cover member functions, including how to write your own, in more detail later.

Also note that `std::string::length()` returns an unsigned integral value (most likely `size_t`). If you want to assign the length to an `int` variable, you should `static_cast` it to avoid compiler warnings about signed/unsigned conversions:

```
1 | int length =  
   | static_cast<int>(myName.length());
```

Conclusion

`std::string` is complex, leveraging many language features that we haven't covered yet. Fortunately, you don't need to understand these complexities to use `std::string` for simple tasks, like basic string input and output. We encourage you to start experimenting with strings now, and we'll cover additional string capabilities later.

Quiz time

Question #1

Write a program that asks the user to enter their full name and their age. As output, tell the user the sum of their age and the number of letters in their name. For simplicity, count spaces in the name as a letter.

Sample output:

```
Enter your full name: John Doe  
Enter your age: 32  
Your age + length of name is: 40
```

Reminder: `std::string::length()` returns an unsigned `int`. You should `static_cast` this to an `int` before adding the age so you don't mix signed and unsigned values.

[Show Solution](#)



Next lesson

4.13 Literals



**Back to table of
contents**



Previous lesson

4.11 Chars

Leave a comment... Put C++ code between triple-backticks (markdown style):````Your C++ code here````



Name*



Email*



Avatars from <https://gravatar.com/> are connected to your provided email address.

Notify me about replies:



POST COMMENT

DP N N FOUT

Newest ▼

