

## 10.7 — An introduction to `std::string_view`

 NASCARDRIVER  AUGUST 25, 2021

In the previous lesson, we talked about C-style strings, and the dangers of using them. C-style strings are fast, but they're not as easy to use and as safe as `std::string`.

But `std::string` (which we covered in [lesson 4.12 -- An introduction to `std::string`](#)), has some of its own downsides, particularly when it comes to const strings.

Consider the following example:

```
1 #include <iostream>
2 #include <string>
3
4 int main()
5 {
6     char text[] { "hello" };
7     std::string str { text };
8     std::string more { str };
9
10    std::cout << text << ' ' << str << ' ' << more <<
11    '\n';
12
13    return 0;
14 }
```

As expected, this prints

```
hello hello hello
```

Internally, `main` copies the string “hello” 3 times, resulting in 4 copies. First, there is the string literal “hello”, which is known at compile-time and stored in the binary. One copy is created when we create the `char[]`. The following two `std::string` objects create one copy of the string each. Because `std::string` is designed to be modifiable, each `std::string` must contain its own copy of the string, so that a given `std::string` can be modified without affecting any other `std::string` object.

This holds true for const `std::string`, even though they can't be modified.

### Introducing `std::string_view`

Consider a window in your house, looking at a car sitting on the street. You can look through the window and see the car, but you can't touch or move the car. Your window just provides a view to the car, which is a completely separate object.

C++17 introduces another way of using strings, `std::string_view`, which lives in the `<string_view>` header.

Unlike `std::string`, which keeps its own copy of the string, `std::string_view` provides a *view* of a string that is defined elsewhere.

We can re-write the above code to use `std::string_view` by replacing every `std::string` with `std::string_view`.

```
1 #include <iostream>
2 #include <string_view>
3
4 int main()
5 {
6     std::string_view text{ "hello" }; // view the text "hello", which is stored in the
7     binary
8     std::string_view str{ text }; // view of the same "hello"
9     std::string_view more{ str }; // view of the same "hello"
10
11     std::cout << text << ' ' << str << ' ' << more << '\n';
12
13     return 0;
14 }
```

The output is the same, but no more copies of the string “hello” are created. The string “hello” is stored in the binary and is not allocated at run-time. `text` is only a view onto the string “hello”, so no copy has to be created. When we copy a `std::string_view`, the new `std::string_view` observes the same string as the copied-from `std::string_view` is observing. This means that neither `str` nor `more` create any copies. They are views onto the existing string “hello”.

`std::string_view` is not only fast, but has many of the functions that we know from `std::string`.

```
1 #include <iostream>
2 #include <string_view>
3
4 int main()
5 {
6     std::string_view str{ "Trains are fast!" };
7
8     std::cout << str.length() << '\n'; // 16
9     std::cout << str.substr(0, str.find(' ')) << '\n'; //
10     Trains
11     std::cout << (str == "Trains are fast!") << '\n'; // 1
12
13     // Since C++20
14     std::cout << str.starts_with("Boats") << '\n'; // 0
15     std::cout << str.ends_with("fast!") << '\n'; // 1
16
17     std::cout << str << '\n'; // Trains are fast!
18
19     return 0;
20 }
```

Because `std::string_view` doesn't create a copy of the string, if we change the viewed string, the changes are reflected in the `std::string_view`.

```

1  #include <iostream>
   #include <string_view>
2
3  int main()
4  {
5      char arr[] { "Gold" };
6      std::string_view str{ arr };
7
8      std::cout << str << '\n'; //
   Gold
9
10     // Change 'd' to 'f' in arr
11     arr[3] = 'f';
12
13     std::cout << str << '\n'; //
   Golf
14
15     return 0;
16 }

```

We modified `arr`, but `str` appears to be changing as well. That's because `arr` and `str` share their string. When you use a `std::string_view`, it's best to avoid modifications to the underlying string for the remainder of the `std::string_view`'s life to prevent confusion and errors.

### Best practice

Use `std::string_view` instead of C-style strings.

Prefer `std::string_view` over `std::string` for read-only strings, unless you already have a `std::string`.

## View modification functions

Back to our window analogy, consider a window with curtains. We can close either the left or right curtain to reduce what we can see. We don't change what's outside, we just reduce the visible area.

Similarly, `std::string_view` contains functions that let us manipulate the *view* of the string. This allows us to change the view without modifying the viewed string.

The functions for this are `remove_prefix`, which removes characters from the left side of the view, and `remove_suffix`, which removes characters from the right side of the view.

```

1 #include <iostream>
2 #include <string_view>

3 int main()
4 {
5     std::string_view str{ "Peach"
6 };

7     std::cout << str << '\n';
8
9     // Ignore the first
10    character.
11    str.remove_prefix(1);

12    std::cout << str << '\n';
13
14    // Ignore the last 2
15    characters.
16    str.remove_suffix(2);

17    std::cout << str << '\n';

18    return 0;
19 }

```

This program produces the following output:

```

Peach
each
ea

```

Unlike real curtains, a `std::string_view` cannot be opened back up. Once you change the visible area, you can't go back (There are tricks which we won't go into).

## std::string\_view works with non-null-terminated strings

Unlike C-style strings and `std::string`, `std::string_view` doesn't use null terminators to mark the end of the string. Rather, it knows where the string ends because it keeps track of its length.

```

1 #include <iostream>
2 #include <iterator> // For std::size
3 #include <string_view>

4 int main()
5 {
6     // No null-terminator.
7     char vowels[] { 'a', 'e', 'i', 'o', 'u' };

8     // vowels isn't null-terminated. We need to pass the length manually.
9     // Because vowels is an array, we can use std::size to get its length.
10    std::string_view str{ vowels, std::size(vowels) };

11    std::cout << str << '\n'; // This is safe. std::cout knows how to print
12    std::string_views.

13    return 0;
14 }

```

This program prints:

```

aeiou

```

## Ownership issues

Being only a view, a `std::string_view`'s lifetime is independent of that of the string it is viewing. If the viewed string goes out of scope, `std::string_view` has nothing to observe and accessing it causes undefined behavior. The string that a `std::string_view` is viewing has to have been created somewhere else. It might be a string literal that lives as long as the program does or it was created by a `std::string`, in which case the string lives until the `std::string` decides to destroy it or the `std::string` dies. `std::string_view` can't create any strings on its own, because it's just a view.

```
What's your name?  
nascardriver  
Hello nascardriver  
Your name is 0P00P0
```

In function `askForName()`, we create `str` and fill it with data from `std::cin`. Then we create `view`, which can view that string. At the end of the function, we return `view`, but the string it is viewing (`str`) is destroyed, so `view` is now pointing to deallocated memory. The function returns a dangling `std::string_view`.

Accessing the returned `std::string_view` in `main` causes undefined behavior, which on the author's machine produced weird characters.

The same can happen when we create a `std::string_view` from a `std::string` and then modify the `std::string`. Modifying a `std::string` can cause its internal string to die and be replaced with a new one in a different place. The `std::string_view` will still look at where the old string was, but it's not there anymore.

## Warning

Make sure that the underlying string viewed with a `std::string_view` does not go out of scope and isn't modified while using the `std::string_view`.

## Converting a `std::string_view` to a `std::string`

An `std::string_view` will not implicitly convert to a `std::string`, but can be explicitly converted:

```

1  #include <iostream>
2  #include <string>
3  #include <string_view>
4
5  void print(const std::string &s)
6  {
7      std::cout << s << '\n';
8  }
9
10 int main()
11 {
12     std::string_view sv{ "balloon" };
13
14     sv.remove_suffix(3);
15
16     // print(sv); // compile error: won't implicitly
convert
17
18     std::string str{ sv }; // okay
19
20     print(str); // okay
21
22     print(static_cast<std::string>(sv)); // okay
23
24     return 0;
25 }

```

This prints:

```

ball
ball

```

## Converting a `std::string_view` to a C-style string

Some old functions (such as the old `strlen` function) still expect C-style strings. To convert a `std::string_view` to a C-style string, we can do so by first converting to a `std::string`:

```

1 #include <cstring>
2 #include <iostream>
3 #include <string>
4 #include <string_view>
5
6 int main()
7 {
8     std::string_view sv{ "balloon" };
9
10    sv.remove_suffix(3);
11
12    // Create a std::string from the std::string_view
13    std::string str{ sv };
14
15    // Get the null-terminated C-style string.
16    const char* szNullTerminated{ str.c_str() };
17
18    // Pass the null-terminated string to the function that we want to use.
19    std::cout << str << " has " << std::strlen(szNullTerminated) << "
20    letter(s)\n";
21
22    return 0;
23 }

```

This prints:

```
ball has 4 letter(s)
```

However, creating a `std::string` every time we want to pass a `std::string_view` as a C-style string is expensive, so this should be avoided if possible.

## Opening the window (kinda) via the data() function

The string being viewed by a `std::string_view` can be accessed by using the `data()` function, which returns a C-style string. This provides fast access to the string being viewed (as a C-string). But it should also only be used if the `std::string_view`'s view hasn't been modified (e.g. by `remove_prefix` or `remove_suffix`) and the string being viewed is null-terminated.

In the following example, `std::strlen` doesn't know what a `std::string_view` is, so we need to pass it `str.data()`:

```

1 #include <cstring> // For std::strlen
2 #include <iostream>
3 #include <string_view>
4
5 int main()
6 {
7     std::string_view str{ "balloon" };
8
9     std::cout << str << '\n';
10
11     // We use std::strlen because it's simple, this could be any other
12     // function
13     // that needs a null-terminated string.
14     // It's okay to use data() because we haven't modified the view, and the
15     // string is null-terminated.
16     std::cout << std::strlen(str.data()) << '\n';
17
18     return 0;
19 }

```

```

balloon
7

```

When a `std::string_view` has been modified, `data()` doesn't always do what we'd like it to. The following example demonstrates what happens when we access `data()` after modifying the view:

```

1 #include <cstring>
2 #include <iostream>
3 #include <string_view>
4
5 int main()
6 {
7     std::string_view str{ "balloon" };
8
9     // Remove the "b"
10    str.remove_prefix(1);
11    // remove the "oon"
12    str.remove_suffix(3);
13    // Remember that the above doesn't modify the string, it only changes
14    // the region that str is observing.
15
16    std::cout << str << " has " << std::strlen(str.data()) << "
17    letter(s)\n";
18    std::cout << "str.data() is " << str.data() << '\n';
19    std::cout << "str is " << str << '\n';
20
21    return 0;
22 }

```

```

all has 6 letter(s)
str.data() is alloon
str is all

```

Clearly this isn't what we'd intended, and is a consequence of trying to access the `data()` of a `std::string_view` that has been modified. The length information about the string is lost when we access `data()`. `std::strlen` and `std::cout` keep reading characters from the underlying string until they find the null-terminator, which is at the end of "balloon".



## Warning

Only use `std::string_view::data()` if the `std::string_view` 's view hasn't been modified and the string being viewed is null-terminated. Using `std::string_view::data()` of a non-null-terminated string can cause undefined behavior.

## Incomplete implementation

Being a relatively recent feature, `std::string_view` isn't implemented as well as it could be.

```
1  std::string s{ "hello" };
2  std::string_view v{ "world" };

3  // Doesn't work
4  std::cout << (s + v) << '\n';
5  std::cout << (v + s) << '\n';

6  // Potentially unsafe, or not what we want, because we're treating
7  // the std::string_view as a C-style string.
8  std::cout << (s + v.data()) << '\n';
9  std::cout << (v.data() + s) << '\n';

10 // Ok, but ugly and wasteful because we have to construct a new
11 // std::string.
12 std::cout << (s + std::string{ v }) << '\n';
13 std::cout << (std::string{ v } + s) << '\n';
14 std::cout << (s + static_cast<std::string>(v)) << '\n';
15 std::cout << (static_cast<std::string>(v) + s) << '\n';
```

There's no reason why line 5 and 6 shouldn't work. They will probably be supported in a future C++ version.



### Next lesson

10.8 Introduction to pointers



### Back to table of contents



### Previous lesson

10.6 C-style strings

Leave a comment... Put C++ code between triple-backticks (markdown style):````Your C++ code````



Name\*



Email\*



Avatars from <https://gravatar.com/> are connected to your provided email address.

Notify me about replies:



POST COMMENT

116 COMMENTS

Newest

