

10.12 — C-style string symbolic constants

ALEX JUNE 17, 2021

C-style string symbolic constants

In a previous lesson, we discussed how you could create and initialize a C-style string, like this:

```
1 #include <iostream>
2 int main()
3 {
4     char myName[] { "Alex" }; // fixed
5     std::cout << myName << '\n';
6     return 0;
7 }
```

C++ also supports a way to create C-style string symbolic constants using pointers:

```
1 #include <iostream>
2 int main()
3 {
4     const char* myName { "Alex" }; // pointer to symbolic
5     std::cout << myName << '\n';
6     return 0;
7 }
```

While these above two programs operate and produce the same results, C++ deals with the memory allocation for these slightly differently.

In the fixed array case, the program allocates memory for a fixed array of length 5, and initializes that memory with the string "Alex\0". Because memory has been specifically allocated for the array, you're free to alter the contents of the array. The array itself is treated as a normal local variable, so when the array goes out of scope, the memory used by the array is freed up for other uses.

In the symbolic constant case, how the compiler handles this is implementation defined. What *usually* happens is that the compiler places the string "Alex\0" into read-only memory somewhere, and then sets the pointer to point to it. Because this memory may be read-only, best practice is to make sure the string is const.

For optimization purposes, multiple string literals may be consolidated into a single value. For example:

```
1 | const char* name1{ "Alex"
   | };
   | const char* name2{ "Alex"
   | };
```

These are two different string literals with the same value. The compiler may opt to combine these into a single shared string literal, with both `name1` and `name2` pointed at the same address. Thus, if `name1` was not `const`, making a change to `name1` could also impact `name2` (which might not be expected).

As a result of string literals being stored in a fixed location in memory, string literals have static duration rather than automatic duration (that is, they die at the end of the program, not the end of the block in which they are defined). That means that when we use string literals, we don't have to worry about scoping issues. Thus, the following is okay:

```
1 | const char*
   | getName()
   | {
2 |     return "Alex";
3 | }
```

In the above code, `getName()` will return a pointer to C-style string "Alex". If this function were returning any other local variable by address, the variable would be destroyed at the end of `getName()`, and we'd return a dangling pointer back to the caller. However, because string literals have static duration, "Alex" will not be destroyed when `getName()` terminates, so the caller can still successfully access it.

C-style strings are used in a lot of old or low-level code, because they have a very small memory footprint. Modern code should favor the use `std::string` and `std::string_view`, as those provide safe and easy access to the string.

`std::cout` and char pointers

At this point, you may have noticed something interesting about the way `std::cout` handles pointers of different types.

Consider the following example:

```
1 | #include <iostream>
2 | int main()
3 | {
4 |     int nArray[5]{ 9, 7, 5, 3, 1 };
5 |     char cArray[]{ "Hello!" };
   |     const char* name{ "Alex" };
6 |
   |     std::cout << nArray << '\n'; // nArray will decay to type
   |     int*
   |     std::cout << cArray << '\n'; // cArray will decay to type
   |     char*
7 |     std::cout << name << '\n'; // name is already type char*
   |
   |     return 0;
   | }
```

On the author's machine, this printed:

```
003AF738
Hello!
Alex
```

Why did the int array print an address, but the character arrays printed strings?

The answer is that `std::cout` makes some assumptions about your intent. If you pass it a non-char pointer, it will simply print the contents of that pointer (the address that the pointer is holding). However, if you pass it an object of type `char*` or `const char*`, it will assume you're intending to print a string. Consequently, instead of printing the pointer's value, it will print the string being pointed to instead!

While this is great 99% of the time, it can lead to unexpected results. Consider the following case:

```
1 #include
  <iostream>
2
3 int main()
4 {
5     char c{ 'Q'
6 };
7     std::cout <<
8     &c;
9
10    return 0;
11 }
```

In this case, the programmer is intending to print the address of variable `c`. However, `&c` has type `char*`, so `std::cout` tries to print this as a string! On the author's machine, this printed:

```
Q|_|_|_|_|4;?A
```

Why did it do this? Well, it assumed `&c` (which has type `char*`) was a string. So it printed the 'Q', and then kept going. Next in memory was a bunch of garbage. Eventually, it ran into some memory holding a 0 value, which it interpreted as a null terminator, so it stopped. What you see may be different depending on what's in memory after variable `c`.

This case is somewhat unlikely to occur in real-life (as you're not likely to actually want to print memory addresses), but it is illustrative of how things work under the hood, and how programs can inadvertently go off the rails.



Next lesson

10.13 Dynamic memory allocation with new and delete



[Back to table of contents](#)



Previous lesson

10.11 Pointer arithmetic and array indexing

Leave a comment... Put C++ code between triple-backticks (markdown style):````Your C++ code here````



Name*



Email*



Avatars from <https://gravatar.com/> are connected to your provided email address.

Notify me about replies:



POST COMMENT

DP N N FOUT

Newest ▼

