

11.13 — Introduction to lambdas (anonymous functions)

👤 NASCARDRIVER 🕒 JULY 15, 2021

Consider this snippet of code that we introduced in [lesson 10.25 -- Introduction to standard library algorithms](#):

```
1 #include <algorithm>
2 #include <array>
3 #include <iostream>
4 #include <string_view>
5
6 static bool containsNut(std::string_view str) // static means internal linkage in this
context
{
    // std::string_view::find returns std::string_view::npos, which is a very large number,
    // if it doesn't find the substring.
    // Otherwise it returns the index where the substring occurs in str.
7     return (str.find("nut") != std::string_view::npos);
8 }
9
10 int main()
11 {
12     constexpr std::array<std::string_view, 4> arr{ "apple", "banana", "walnut", "lemon" };
13
14     // std::find_if takes a pointer to a function
15     const auto found{ std::find_if(arr.begin(), arr.end(), containsNut) };
16
17     if (found == arr.end())
18     {
19         std::cout << "No nuts\n";
20     }
21     else
22     {
23         std::cout << "Found " << *found << '\n';
24     }
25
26     return 0;
27 }
```

This code searches through an array of strings looking for the first element that contains the substring “nut”. Thus, it produces the result:

```
Found walnut
```

And while it works, it could be improved.

The root of the issue here is that `std::find_if` requires that we pass it a function pointer. Because of that, we are forced to define a function that’s only going to be used once, that must be given a name, and that must be put in the global scope (because functions can’t be nested!). The function is also so short, it’s almost easier to discern what it does from the one line of code than from the name and comments.

Lambdas to the rescue

A lambda expression (also called a lambda or closure) allows us to define an anonymous function inside another function. The nesting is important, as it allows us both to avoid namespace naming pollution, and to define the function as close to where it is used as possible (providing additional context).

The syntax for lambdas is one of the weirder things in C++, and takes a bit of getting used to. Lambdas take the form:

```
[ captureClause ] ( parameters ) -> returnType
{
    statements;
}
```

The `capture clause` and `parameters` can both be empty if they are not needed.

The `return type` is optional, and if omitted, `auto` will be assumed (thus using type inference used to determine the return type). While we previously noted that type inference for function return types should be avoided, in this context, it's fine to use (because these functions are typically so trivial).

Also note that lambdas have no name, so we don't need to provide one.

As an aside...

This means a trivial lambda definition looks like this:

```
1  #include <iostream>
2  int main()
3  {
4      [] {}; // defines a lambda with no captures, no parameters, and no return
5      type
           return 0;
   }
```

Let's rewrite the above example using a lambda:

```

1 #include <algorithm>
2 #include <array>
3 #include <iostream>
4 #include <string_view>
5
6 int main()
7 {
8     constexpr std::array<std::string_view, 4> arr{ "apple", "banana", "walnut", "lemon" };
9
10    // Define the function right where we use it.
11    const auto found{ std::find_if(arr.begin(), arr.end(),
12                                [](std::string_view str) // here's our lambda, no capture
13                                {
14                                    return (str.find("nut") != std::string_view::npos);
15                                }) };
16
17    if (found == arr.end())
18    {
19        std::cout << "No nuts\n";
20    }
21    else
22    {
23        std::cout << "Found " << *found << '\n';
24    }
25
26    return 0;
27 }

```

This works just like the function pointer case, and produces an identical result:

```
Found walnut
```

Note how similar our lambda is to our `containsNut` function. They both have identical parameters and function bodies. The lambda has no capture clause (we'll explain what a capture clause is in the next lesson) because it doesn't need one. And we've omitted the trailing return type in the lambda (for conciseness), but since `operator!=` returns a `bool`, our lambda will return a `bool` too.

Type of a lambda

In the above example, we defined a lambda right where it was needed. This use of a lambda is sometimes called a function literal.

However, writing a lambda in the same line as it's used can sometimes make code harder to read. Much like we can initialize a variable with a literal value (or a function pointer) for use later, we can also initialize a lambda variable with a lambda definition and then use it later. A named lambda along with a good function name can make code easier to read.

For example, in the following snippet, we're using `std::all_of` to check if all elements of an array are even:

```

1 // Bad: We have to read the lambda to understand what's happening.
2 return std::all_of(array.begin(), array.end(), [](int i){ return ((i % 2) == 0);
3 });

```

We can improve the readability of this as follows:

```

1 // Good: Instead, we can store the lambda in a named variable and pass it to the
  function.
  auto isEven{
    [](int i)
    {
2      return ((i % 2) == 0);
3    }
4  };
5
6
7 return std::all_of(array.begin(), array.end(), isEven);

```

Note how well the last line reads: “return whether *all* of the elements in the *array* are *even*”

But what is the type of lambda `isEven`?

As it turns out, lambdas don’t have a type that we can explicitly use. When we write a lambda, the compiler generates a unique type just for the lambda that is not exposed to us.

For advanced readers

In actuality, lambdas aren’t functions (which is part of how they avoid the limitation of C++ not supporting nested functions). They’re a special kind of object called a functor. Functors are objects that contain an overloaded `operator()` that make them callable like a function.

Although we don’t know the type of a lambda, there are several ways of storing a lambda for use post-definition. If the lambda has an empty capture clause, we can use a regular function pointer. In the next lesson, we introduce lambda captures, a function pointer won’t work anymore at that point. However, `std::function` can be used for lambdas even if they are capturing something.

```

1 #include <functional>
2
3 int main()
4 {
5     // A regular function pointer. Only works with an empty capture clause.
    double (*addNumbers1)(double, double){
6         [](double a, double b) {
7             return (a + b);
8         }
9     };
10
11     addNumbers1(1, 2);
12
13     // Using std::function. The lambda could have a non-empty capture clause (Next lesson).
14     std::function addNumbers2{ // note: pre-C++17, use std::function<double(double, double)>
15         instead
16         [](double a, double b) {
17             return (a + b);
18         }
19     };
20
21     addNumbers2(3, 4);
22
23     // Using auto. Stores the lambda with its real type.
    auto addNumbers3{
24         [](double a, double b) {
25             return (a + b);
26         }
27     };
28
29     addNumbers3(5, 6);
30
31     return 0;
32 }

```

The only way of using the lambda’s actual type is by means of `auto`. `auto` also has the benefit of having no overhead compared to `std::function`.

Unfortunately, we can’t always use `auto`. In cases where the actual lambda is unknown (e.g. because we’re passing a lambda to a function as a parameter and the caller determines what lambda will be passed in), we can’t use `auto` without compromises. In such cases, `std::function` can be used.

```

1 #include <functional>
2 #include <iostream>
3
4 // We don't know what fn will be. std::function works with regular functions and
  lambdas.
  void repeat(int repetitions, const std::function<void(int)>& fn)
  {
5     for (int i{ 0 }; i < repetitions; ++i)
6     {
7         fn(i);
8     }
9 }
10
11 int main()
12 {
13     repeat(3, [](int i) {
14         std::cout << i << '\n';
15     });
16 }

```

Output

```

0
1
2

```

If we had used `auto` for the type of `fn`, the caller of the function wouldn't know what parameters and return type `fn` needs to have. Furthermore, functions with `auto` parameters cannot be separated into a header and source file. We cover the reason for this restriction when we talk about templates.

Rule

Use `auto` when initializing variables with lambdas, and `std::function` if you can't initialize the variable with the lambda.

Generic lambdas

For the most part, lambda parameters work by the same rules as regular function parameters.

One notable exception is that since C++14 we're allowed to use `auto` for parameters (note: in C++20, regular functions are able to use `auto` for parameters too). When a lambda has one or more `auto` parameter, the compiler will infer what parameter types are needed from the calls to the lambda.

Because lambdas with one or more `auto` parameter can potentially work with a wide variety of types, they are called generic lambdas.

For advanced readers

When used in the context of a lambda, `auto` is just a shorthand for a template parameter.

Let's take a look at a generic lambda:

```

1 #include <algorithm>
2 #include <array>
3 #include <iostream>
4 #include <string_view>
5
6 int main()
7 {
8     constexpr std::array months{ // pre-C++17 use std::array<const char*,
12>
9         "January",
10        "February",
11        "March",
12        "April",
13        "May",
14        "June",
15        "July",
16        "August",
17        "September",
18        "October",
19        "November",
20        "December"
21    };
22
23    // Search for two consecutive months that start with the same letter.
24    const auto sameLetter{ std::adjacent_find(months.begin(), months.end(),
25        [](const auto& a, const auto& b) {
26            return (a[0] == b[0]);
27        }) };
28
29    // Make sure that two months were found.
30    if (sameLetter != months.end())
31    {
32        // std::next returns the next iterator after sameLetter
33        std::cout << *sameLetter << " and " << *std::next(sameLetter)
34            << " start with the same letter\n";
35    }
36
37    return 0;
38 }

```

Output:

```
June and July start with the same letter
```

In the above example, we use `auto` parameters to capture our strings by `const` reference. Because all string types allow access to their individual characters via `operator[]`, we don't need to care whether the user is passing in a `std::string`, C-style string, or something else. This allows us to write a lambda that could accept any of these, meaning if we change the type of `months` later, we won't have to rewrite the lambda.

However, `auto` isn't always the best choice. Consider:

```

1  #include <algorithm>
2  #include <array>
3  #include <iostream>
4  #include <string_view>
5
6  int main()
7  {
8      constexpr std::array months{ // pre-C++17 use std::array<const char*, 12>
9          "January",
10         "February",
11         "March",
12         "April",
13         "May",
14         "June",
15         "July",
16         "August",
17         "September",
18         "October",
19         "November",
20         "December"
21     };
22
23     // Count how many months consist of 5 letters
24     const auto fiveLetterMonths{ std::count_if(months.begin(), months.end(),
25         [](std::string_view str) {
26             return (str.length() == 5);
27         }) };
28
29     std::cout << "There are " << fiveLetterMonths << " months with 5
30     letters\n";
31
32     return 0;
33 }

```

Output:

```
There are 2 months with 5 letters
```

In this example, using `auto` would infer a type of `const char*`. C-style strings aren't easy to work with (apart from using `operator[]`). In this case, we prefer to explicitly define the parameter as a `std::string_view`, which allows us to work with the underlying data much more easily (e.g. we can ask the string view for its length, even if the user passed in a C-style array).

Generic lambdas and static variables

One thing to be aware of is that a unique lambda will be generated for each different type that `auto` resolves to. The following example shows how one generic lambda turns into two distinct lambdas:

```

1  #include <algorithm>
   #include <array>
2  #include <iostream>
   #include <string_view>
3
4  int main()
5  {
6      // Print a value and count how many times @print has been
7      // called.
8      auto print{
9          [](auto value) {
10             static int callCount{ 0 };
11             std::cout << callCount++ << ": " << value << '\n';
12         }
13     };
14
15     print("hello"); // 0: hello
16     print("world"); // 1: world
17
18     print(1); // 0: 1
19     print(2); // 1: 2
20
21     print("ding dong"); // 2: ding dong
22
23     return 0;
24 }

```

Output

```

0: hello
1: world
0: 1
1: 2
2: ding dong

```

In the above example, we define a lambda and then call it with two different parameters (a string literal parameter, and an integer parameter). This generates two different versions of the lambda (one with a string literal parameter, and one with an integer parameter).

Most of the time, this is inconsequential. However, note that if the generic lambda uses static duration variables, those variables are not shared between the generated lambdas.

We can see this in the example above, where each type (string literals and integers) has its own unique count! Although we only wrote

the lambda once, two lambdas were generated -- and each has its own version of `callCount`. To have a shared counter between the two generated lambdas, we'd have to define a global variable or a `static` local variable outside of the lambda. As you know from previous lessons, both global- and static local variables can cause problems and make it more difficult to understand code. We'll be able to avoid those variables after talking about lambda captures in the next lesson.

Return type deduction and trailing return types

If return type deduction is used, a lambda's return type is deduced from the `return`-statements inside the lambda. If return type inference is used, all return statements in the lambda must return the same type (otherwise the compiler won't know which one to prefer).

For example:

```
1 #include <iostream>
2
3 int main()
4 {
5     auto divide{ [](int x, int y, bool bInteger) { // note: no specified return type
6         if (bInteger)
7             return x / y;
8         else
9             return static_cast<double>(x) / y; // ERROR: return type doesn't match previous return
10    } };
11
12    std::cout << divide(3, 2, true) << '\n';
13    std::cout << divide(3, 2, false) << '\n';
14
15    return 0;
16 }
```

This produces a compile error because the return type of the first return statement (`int`) doesn't match the return type of the second return statement (`double`).

In the case where we're returning different types, we have two options:

1. Do explicit casts to make all the return types match, or
2. explicitly specify a return type for the lambda, and let the compiler do implicit conversions.

The second case is usually the better choice:

```
1 #include <iostream>
2
3 int main()
4 {
5     // note: explicitly specifying this returns a double
6     auto divide{ [](int x, int y, bool bInteger) -> double {
7         if (bInteger)
8             return x / y; // will do an implicit conversion to
9         double
10        else
11            return static_cast<double>(x) / y;
12    } };
13
14    std::cout << divide(3, 2, true) << '\n';
15    std::cout << divide(3, 2, false) << '\n';
16
17    return 0;
18 }
```

That way, if you ever decide to change the return type, you (usually) only need to change the lambda's return type, and not touch the lambda body.

Standard library function objects

For common operations (e.g. addition, negation, or comparison) you don't need to write your own lambdas, because the standard library comes with many basic callable objects that can be used instead. These are defined in the `<functional>` header.

In the following example:

```

1 #include <algorithm>
2 #include <array>
3 #include <iostream>
4
5 bool greater(int a, int b)
6 {
7     // Order @a before @b if @a is greater than
8     @b.
9     return (a > b);
10 }
11
12 int main()
13 {
14     std::array arr{ 13, 90, 99, 5, 40, 80 };
15
16     // Pass greater to std::sort
17     std::sort(arr.begin(), arr.end(), greater);
18
19     for (int i : arr)
20     {
21         std::cout << i << ' ';
22     }
23
24     std::cout << '\n';
25
26     return 0;
27 }

```

Output

```
99 90 80 40 13 5
```

Instead of converting our `greater` function to a lambda (which would obscure its meaning a bit), we can instead use `std::greater :`

```

1 #include <algorithm>
2 #include <array>
3 #include <iostream>
4 #include <functional> // for std::greater
5
6 int main()
7 {
8     std::array arr{ 13, 90, 99, 5, 40, 80 };
9
10    // Pass std::greater to std::sort
11    std::sort(arr.begin(), arr.end(), std::greater{}); // note: need curly braces to instantiate
12    object
13
14    for (int i : arr)
15    {
16        std::cout << i << ' ';
17    }
18
19    std::cout << '\n';
20
21    return 0;
22 }

```

Output

```
99 90 80 40 13 5
```

Conclusion

Lambdas and the algorithm library may seem unnecessarily complicated when compared to a solution that uses a loop. However, this combination can allow some very powerful operations in just a few lines of code, and can be more readable than writing your own loops. On top of that, the algorithm library features powerful and easy-to-use parallelism, which you won't get with loops. Upgrading source code that uses library functions is easier than upgrading code that uses loops.

Lambdas are great, but they don't replace regular functions for all cases. Prefer regular functions for non-trivial and reusable cases.

Quiz time

Question #1

Create a `struct Student` that stores the name and points of a student. Create an array of students and use `std::max_element` to find the student with the most points, then print that student's name. `std::max_element` takes the `begin` and `end` of a list, and a function that takes 2 parameters and returns `true` if the first argument is less than the second.

Given the following array

```
1  std::array<Student, 8> arr{
2      { "Albert", 3 },
3      { "Ben", 5 },
4      { "Christine", 2 },
5      { "Dan", 8 }, // Dan has the most points
6  };
```

your program should print

```
Dan is the best student
```

[Show Hint](#)

[Show Solution](#)

Question #2

Use `std::sort` and a lambda in the following code to sort the seasons by ascending average temperature.

```
1  #include <algorithm>
2  #include <array>
3  #include <iostream>
4  #include <string_view>
5
6  struct Season
7  {
8      std::string_view name{};
9      double averageTemperature{};
10 };
11
12 int main()
13 {
14     std::array<Season, 4> seasons{
15         { "Spring", 285.0 },
16         { "Summer", 296.0 },
17         { "Fall", 288.0 },
18         { "Winter", 263.0 }
19     };
20
21     /*
22      * Use std::sort here
23      */
24
25     for (const auto& season :
26 seasons)
27     {
28         std::cout << season.name <<
29 '\n';
30     }
31
32     return 0;
33 }
```

The program should print

Winter
Spring
Fall
Summer

Show Solution



Next lesson

11.14 Lambda captures



Back to table of contents



Previous lesson

11.12 Ellipsis (and why to avoid them)

Leave a comment...``Place code to highlight between triple-backticks (markdown style)``



Name*



Email*



Avatars from <https://gravatar.com/> are connected to your provided email address.

Notify me about replies:



POST COMMENT

