

## 18.3 — The override and final specifiers, and covariant return types

ALEX SEPTEMBER 28, 2021

To address some common challenges with inheritance, there are two special identifiers: **override** and **final**. Note that these identifiers are not considered keywords -- they are normal identifiers that have special meaning in certain contexts.

Although **final** isn't used very much, **override** is a fantastic addition that you should use regularly. In this lesson, we'll take a look at both, as well as one exception to the rule that virtual function override return types must match.

### The override specifier

As we mentioned in the previous lesson, a derived class virtual function is only considered an override if its signature and return types match exactly. That can lead to inadvertent issues, where a function that was intended to be an override actually isn't.

Consider the following example:

```
1 #include <iostream>
2 #include <string_view>
3
4 class A
5 {
6 public:
7     virtual std::string_view getName1(int x) { return "A"; }
8     virtual std::string_view getName2(int x) { return "A"; }
9 };
10
11 class B : public A
12 {
13 public:
14     virtual std::string_view getName1(short int x) { return "B"; } // note: parameter is a short
15     virtual std::string_view getName2(int x) const { return "B"; } // note: function is const
16 };
17
18 int main()
19 {
20     B b{};
21     A& rBase{ b };
22     std::cout << rBase.getName1(1) << '\n';
23     std::cout << rBase.getName2(2) << '\n';
24
25     return 0;
26 }
```

Because `rBase` is an `A` reference to a `B` object, the intention here is to use virtual functions to access `B::getName1()` and `B::getName2()`. However, because `B::getName1()` takes a different parameter (a `short int` instead of an `int`), it's not considered an override of `A::getName1()`. More insidiously, because `B::getName2()` is `const` and `A::getName2()` isn't, `B::getName2()` isn't considered an override of `A::getName2()`.

Consequently, this program prints:

```
A
A
```

In this particular case, because `A` and `B` just print their names, it's fairly easy to see that we messed up our overrides, and that the wrong virtual function is being called. However, in a more complicated program, where the functions have behaviors or return values that aren't printed, such issues can be very difficult to debug.

To help address the issue of functions that are meant to be overrides but aren't, the override specifier can be applied to any virtual function by placing the specifier in the same place const would go. If the function does not override a base class function (or is applied to a non-virtual function), the compiler will flag the function as an error.

```
1 #include <string_view>
2
3 class A
4 {
5 public:
6     virtual std::string_view getName1(int x) { return "A"; }
7     virtual std::string_view getName2(int x) { return "A"; }
8     virtual std::string_view getName3(int x) { return "A"; }
9 };
10
11 class B : public A
12 {
13 public:
14     std::string_view getName1(short int x) override { return "B"; } // compile error, function is not an
15     override
16     std::string_view getName2(int x) const override { return "B"; } // compile error, function is not an
17     override
18     std::string_view getName3(int x) override { return "B"; } // okay, function is an override of
19     A::getName3(int)
20 };
21
22 int main()
23 {
24     return 0;
25 }
```

The above program produces two compile errors: one for B::getName1(), and one for B::getName2(), because neither override a prior function. B::getName3() does override A::getName3(), so no error is produced for that line.

Because there is no performance penalty for using the override specifier and it helps ensure you've actually overridden the function you think you have, all virtual override functions should be tagged using the override specifier. Additionally, because the override specifier implies virtual, there's no need to tag functions using the override specifier with the virtual keyword.

### Best practice

Use the virtual keyword on virtual functions in a base class. Use the override specifier (but not the virtual keyword) on override functions in derived classes.

### The final specifier

There may be cases where you don't want someone to be able to override a virtual function, or inherit from a class. The final specifier can be used to tell the compiler to enforce this. If the user tries to override a function or inherit from a class that has been specified as final, the compiler will give a compile error.

In the case where we want to restrict the user from overriding a function, the final specifier is used in the same place the override specifier is, like so:

```

1  #include <string_view>
2  class A
3  {
4  public:
5      virtual std::string_view getName() { return "A"; }
6  };
7
8      class B : public A
9      {
10     public:
11         // note use of final specifier on following line -- that makes this function no longer overridable
12         std::string_view getName() override final { return "B"; } // okay, overrides A::getName()
13     };
14
15     class C : public B
16     {
17     public:
18         std::string_view getName() override { return "C"; } // compile error: overrides B::getName(), which is
19         final
20     };

```

In the above code, B::getName() overrides A::getName(), which is fine. But B::getName() has the final specifier, which means that any further overrides of that function should be considered an error. And indeed, C::getName() tries to override B::getName() (the override specifier here isn't relevant, it's just there for good practice), so the compiler will give a compile error.

In the case where we want to prevent inheriting from a class, the final specifier is applied after the class name:

```

1  #include <string_view>
2  class A
3  {
4  public:
5      virtual std::string_view getName() { return "A"; }
6  };
7
8      class B final : public A // note use of final specifier here
9      {
10     public:
11         std::string_view getName() override { return "B"; }
12     };
13
14     class C : public B // compile error: cannot inherit from final
15     class
16     {
17     public:
18         std::string_view getName() override { return "C"; }
19     };

```

In the above example, class B is declared final. Thus, when C tries to inherit from B, the compiler will give a compile error.

### Covariant return types

There is one special case in which a derived class virtual function override can have a different return type than the base class and still be considered a matching override. If the return type of a virtual function is a pointer or a reference to a class, override functions can return a pointer or a reference to a derived class. These are called covariant return types. Here is an example:

```

1  #include <iostream>
2  #include <string_view>

3  class Base
4  {
5  public:
6      // This version of getThis() returns a pointer to a Base class
7      virtual Base* getThis() { std::cout << "called Base::getThis()\n"; return this; }
8      void printType() { std::cout << "returned a Base\n"; }
9  };

10 class Derived : public Base
11 {
12 public:
13     // Normally override functions have to return objects of the same type as the base function
14     // However, because Derived is derived from Base, it's okay to return Derived* instead of Base*
15     Derived* getThis() override { std::cout << "called Derived::getThis()\n"; return this; }
16     void printType() { std::cout << "returned a Derived\n"; }
17 };

18 int main()
19 {
20     Derived d{};
21     Base* b{ &d };
22     d.getThis()->printType(); // calls Derived::getThis(), returns a Derived*, calls
23     Derived::printType
24     b->getThis()->printType(); // calls Derived::getThis(), returns a Base*, calls Base::printType

25     return 0;
26 }

```

This prints:

```

called Derived::getThis()
returned a Derived
called Derived::getThis()
returned a Base

```

One interesting note about covariant return types: C++ can't dynamically select types, so you'll always get the type that matches the actual version of the function being called.

In the above example, we first call `d.getThis()`. Since `d` is a `Derived`, this calls `Derived::getThis()`, which returns a `Derived*`. This `Derived*` is then used to call non-virtual function `Derived::printType()`.

Now the interesting case. We then call `b->getThis()`. Variable `b` is a `Base` pointer to a `Derived` object. `Base::getThis()` is a virtual function, so this calls `Derived::getThis()`. Although `Derived::getThis()` returns a `Derived*`, because `Base` version of the function returns a `Base*`, the returned `Derived*` is upcast to a `Base*`. Because `Base::printType()` is non-virtual, `Base::printType()` is called.

In other words, in the above example, you only get a `Derived*` if you call `getThis()` with an object that is typed as a `Derived` object in the first place.

Note that if `printType()` were virtual instead of non-virtual, the result of `b->getThis()` (an object of type `Base*`) would have undergone virtual function resolution, and `Derived::printType()` would have been called.



## Next lesson

**18.4** Virtual destructors, virtual assignment, and overriding virtualization



Back to table of contents



## Previous lesson

**18.2** Virtual functions and polymorphism

Leave a comment...``Place code to highlight between triple-backticks (markdown style)``

 Name\*

 Email\* 

Avatars from <https://gravatar.com/> are connected to your provided email address.

Notify me about replies:



POST COMMENT

81 COMMENTS

Newest ▼

