

M.2 — R-value references

ALEX AUGUST 2, 2021

Way back in chapter 1, we mentioned l-values and r-values, and then told you not to worry that much about them. That was fair advice prior to C++11. But understanding move semantics in C++11 requires a re-examination of the topic. So let's do that now.

L-values and r-values

Despite having the word “value” in their names, l-values and r-values are actually not properties of values, but rather, properties of expressions.

Every expression in C++ has two properties: a type (which is used for type checking), and a value category (which is used for certain kinds of syntax checking, such as whether the result of the expression can be assigned to). In C++03 and earlier, l-values and r-values were the only two value categories available.

The actual definition of which expressions are l-values and which are r-values is surprisingly complicated, so we'll take a simplified view of the subject that will

largely suffice for our purposes.

It's simplest to think of an l-value (also called a locator value) as a function or an object (or an expression that evaluates to a function or object). All l-values have assigned memory addresses.

When l-values were originally defined, they were defined as “values that are suitable to be on the left-hand side of an assignment expression”. However, later, the `const` keyword was added to the language, and l-values were split into two sub-categories: modifiable l-values, which can be changed, and non-modifiable l-values, which are `const`.

It's simplest to think of an r-value as “everything that is not an l-value”. This notably includes literals (e.g. 5), temporary values (e.g. `x+1`), and anonymous objects (e.g. `Fraction(5, 2)`). r-values are typically evaluated for their values, have expression scope (they die at the end of the expression they are in), and cannot be assigned to. This non-assignment rule makes sense, because assigning a value applies a side-effect to the object. Since r-values have expression scope, if we were to assign a value to an r-value, then the r-value would either go out of scope before we had a chance to use the assigned value in the next expression (which makes the assignment useless) or we'd have to use a variable with a side effect applied more than once in an expression (which by now you should know causes undefined behavior!).

In order to support move semantics, C++11 introduces 3 new value categories: pr-values, x-values, and gl-values. We will largely ignore these since understanding them isn't necessary to learn about or use move semantics effectively. If you're interested, cppreference.com has an extensive list of expressions that qualify for each of the various value categories, as well as more detail about them.

L-value references

Prior to C++11, only one type of reference existed in C++, and so it was just called a “reference”. However, in C++11, it's sometimes called an l-value reference. L-value references can only be initialized with modifiable l-values.

L-value reference	Can be initialized with	Can modify
Modifiable l-values	Yes	Yes
Non-modifiable l-values	No	No
R-values	No	No

L-value references to const objects can be initialized with l-values and r-values alike. However, those values can't be modified.

L-value reference to const	Can be initialized with	Can modify
Modifiable l-values	Yes	No
Non-modifiable l-values	Yes	No
R-values	Yes	No

L-value references to const objects are particularly useful because they allow us to pass any type of argument (l-value or r-value) into a function without making a copy of the argument.

R-value references

C++11 adds a new type of reference called an r-value reference. An r-value reference is a reference that is designed to be initialized with an r-value (only). While an l-value reference is created using a single ampersand, an r-value reference is created using a double ampersand:

```

1 | int x{ 5 };
  | int &lref{ x }; // l-value reference initialized with l-value
2 | x
  | int &&rref{ 5 }; // r-value reference initialized with r-
  | value 5

```

R-values references cannot be initialized with l-values.

R-value reference	Can be initialized with	Can modify
Modifiable l-values	No	No
Non-modifiable l-values	No	No
R-values	Yes	Yes

R-value reference to const	Can be initialized with	Can modify
Modifiable l-values	No	No
Non-modifiable l-values	No	No
R-values	Yes	No

R-value references have two properties that are useful. First, r-value references extend the lifespan of the object they are initialized with to the lifespan of the r-value reference (l-value references to const objects can do this too). Second, non-const r-value references allow you to modify the r-value!

Let's take a look at some examples:

```

1 #include <iostream>
2
3 class Fraction
4 {
5 private:
6     int m_numerator;
7     int m_denominator;
8
9 public:
10    Fraction(int numerator = 0, int denominator = 1) :
11        m_numerator{ numerator }, m_denominator{ denominator }
12    {
13
14        friend std::ostream& operator<<(std::ostream& out, const Fraction &f1)
15        {
16            out << f1.m_numerator << '/' << f1.m_denominator;
17            return out;
18        }
19    };
20
21 int main()
22 {
23     auto &rref{ Fraction{ 3, 5 } }; // r-value reference to temporary
24     Fraction
25
26         // f1 of operator<< binds to the temporary, no copies are created.
27         std::cout << rref << '\n';
28
29     return 0;
30 } // rref (and the temporary Fraction) goes out of scope here

```

This program prints:

```
3/5
```

As an anonymous object, `Fraction(3, 5)` would normally go out of scope at the end of the expression in which it is defined. However, since we're initializing an r-value reference with it, its duration is extended until the end of the block. We can then use that r-value reference to print the Fraction's value.

Now let's take a look at a less intuitive example:

```

1 #include <iostream>
2
3 int main()
4 {
5     int &rref{ 5 }; // because we're initializing an r-value reference with a literal, a temporary with
6     value 5 is created here
7     rref = 10;
8     std::cout << rref << '\n';
9
10    return 0;
11 }

```

This program prints:

```
10
```

While it may seem weird to initialize an r-value reference with a literal value and then be able to change that value, when initializing an r-value with a literal, a temporary is constructed from the literal so that the reference is referencing a temporary object, not a literal value.

R-value references are not very often used in either of the manners illustrated above.

R-value references as function parameters

R-value references are more often used as function parameters. This is most useful for function overloads when you want to have different behavior for l-value and r-value arguments.

```
1 void fun(const int &lref) // l-value arguments will select this
  function
  {
    std::cout << "l-value reference to const\n";
  }

  void fun(int &&rref) // r-value arguments will select this function
  {
2   std::cout << "r-value reference\n";
3  }

  int main()
  {
4   int x{ 5 };
5   fun(x); // l-value argument calls l-value version of function
6   fun(5); // r-value argument calls r-value version of function

   return 0;
  }
```

This prints:

```
l-value reference to const
r-value reference
```

As you can see, when passed an l-value, the overloaded function resolved to the version with the l-value reference. When passed an r-value, the overloaded function resolved to the version with the r-value reference (this is considered a better match than a l-value reference to const).

Why would you ever want to do this? We'll discuss this in more detail in the next lesson. Needless to say, it's an important part of move semantics.

One interesting note:

```
1 int &&ref{ 5
  };
  fun(ref);
```

actually calls the l-value version of the function! Although variable `ref` has type *r-value reference to an integer*, it is actually an l-value itself (as are all named variables). The confusion stems from the use of the term r-value in two different contexts. Think of it this way: Named-objects are l-values. Anonymous objects are r-values. The type of the named object or anonymous object is independent from whether it's an l-value or r-value. Or, put another way, if r-value reference had been called anything else, this confusion wouldn't exist.

Returning an r-value reference

You should almost never return an r-value reference, for the same reason you should almost never return an l-value reference. In most cases, you'll end up returning a hanging reference when the referenced object goes out of scope at the end of the function.

Quiz time

1. State which of the following lettered statements will not compile:

```
1  int main()
2  {
3      int x{};
4
5      // l-value references
6      int &ref1{ x }; // A
7      int &ref2{ 5 }; // B
8
9      const int &ref3{ x }; //
10     C
11     const int &ref4{ 5 }; //
12     D
13
14     // r-value references
15     int &&ref5{ x }; // E
16     int &&ref6{ 5 }; // F
17
18     const int &&ref7{ x };
19     // G
20     const int &&ref8{ 5 };
21     // H
22
23     return 0;
24 }
```

Show Solution



Next lesson

M.3 Move constructors and move assignment



Back to table of contents



Previous lesson

M.1 Intro to smart pointers and move semantics

Leave a comment... Put C++ code between triple-backticks (markdown style):````Your C++ code`



Name*



Email*



Avatars from <https://gravatar.com/> are connected to your provided email address.

Notify me about replies:



POST COMMENT

DP N N FOUT

Newest ▼

